

Lab 1: CMSIS-DSP Functions for Optimized Code

ECSE 444: Microprocessors

Winter 2019

Introduction

In this lab you will be introduced to the [CMSIS-DSP software library](#), which is a collection of common signal processing functions that ARM has optimized for the various Cortex-M processor cores.

In particular, you will:

- Write two simple tasks in C code
- Re-write the tasks using the CMSIS-DSP library functions.
- Learn how to time the code to evaluate the speed-up
- Write the tasks in assembly to evaluate the speed-up (or lack thereof) of hand written assembly v/s C.

Changelog

- 17-Jan Initial release.

Grading

- 25 % C code for dot-product
- 10 % CMSIS-DSP implementation of dot product
- 5 % Timing and analysis of both implementations
- 25 % C code for variance
- 10 % CMSIS-DSP implementation of variance
- 5 % Timing and analysis of both implementations
- 10 % Assembly code for dot-product
- 10 % Assembly code for variance

Note: *Please remember that your assembly code must run faster than your C code!*

Background

This section will introduce you to all the concepts you need to succeed in this lab.

As an example, we will provide the code to find the maximum value in an array of floating point elements. You are provided with a base-project **GXX_Lab1** which contains the files [main.c](#), [asm_math.s](#) and [asm_math.h](#).

IMPORTANT: If you are using the lab computers, please ensure that you save the project in C:\Users\<your_name_goes_here> otherwise it will not work in Keil.

```

1  #include <stdio.h>
2  #include <arm_math.h>
3
4  #include <asm_max.h>
5
6  float fl0_array[10] = {48.21, 79.48, 24.27, 28.82, 78.24, 88.49, 31.19, 5.52, 82.70, 77.73};
7
8  int main() {
9
10     float max;
11     int max_idx;
12
13     float a_max;
14     int a_max_idx;
15
16     float DSP_max;
17     uint32_t DSP_max_idx;
18
19     int i;
20
21     /* LOOP TO FIND MAX*/
22     max = fl0_array[0];
23     max_idx = 0;
24     for(i=0 ; i<10 ; i++) {
25         if(fl0_array[i] > max) {
26             max = fl0_array[i];
27             max_idx = i;
28         }
29     }
30
31     /* Assembly max function */
32     asm_max(fl0_array, 10, &a_max, &a_max_idx);
33
34     /* CMSIS-DSP max function */
35     arm_max_f32(fl0_array, 10, &DSP_max, &DSP_max_idx);
36
37     printf("Pure C   : Max element f[%d] = %f\n", max_idx, max);
38     printf("Assembly : Max element f[%d] = %f\n", a_max_idx, a_max);
39     printf("C: DSP    : Max element f[%d] = %f\n", DSP_max_idx, DSP_max);
40     return 0;
41 }
42
```

Figure 1: main.c

```

1  #ifndef _ASM_MAX
2  #define _ASM_MAX
3
4     void asm_max(float *f_arr, int N, float *max, int *max_idx);
5
6  #endif
7

```

Figure 2: asm_max.h

```

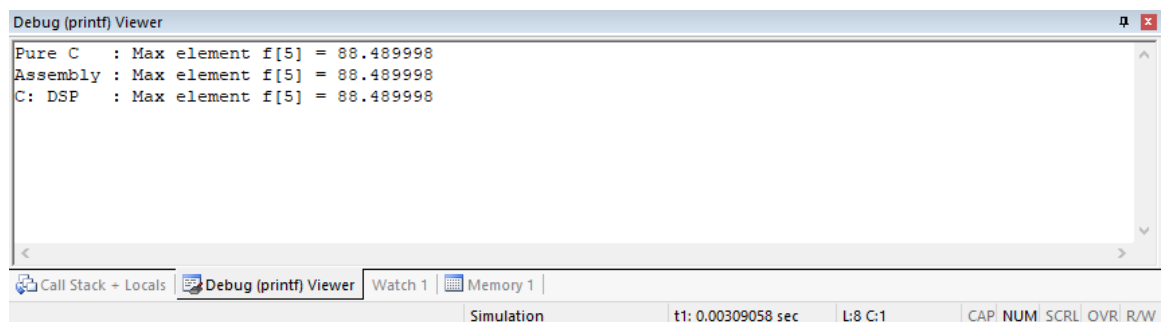
1      AREA test, CODE, READONLY
2
3      export asm_max                ; label must be exported if it is to be used as a function in C
4      asm_max
5
6      PUSH{R4, R5}                  ; saving context according to calling convention
7
8      VLDR.f32 S0, [R0]              ; max = f[0] (register S0 used for max)
9      MOV R4, #0                    ; max_index = 0 (register R4 used for max_idx)
10
11  loop
12      SUBS R1, R1, #1                ; loop counter (N = N-1)
13      BLT done                      ; loop has finished?
14      ADD R5, R0, R1, LSL #2        ; creating base address for the next element in the array
15      VLDR.f32 S1, [R5]             ; load next element into S1
16      VCMPEQ.f32 S0, S1              ; compare element with current max
17      VMRS APSR_nzcv, FPSCR         ; need to move the FPU status register to achieve floating point conditional execution
18      BGT continue                  ; if the current max is greater or equal to new element, no need to update, continue
19      VMOV.f32 S0, S1               ; if not, then update current max
20      MOV R4, R1                    ; also update max index
21
22  continue
23      B loop                        ; loop
24
25  done
26      VSTR.f32 S0, [R2]              ; store the max value in the pointer (float *max) that was provided
27      STR R4, [R3]                  ; store the index of the max value in the pointer (int *max_idx) that was provided
28
29      POP{R4, R5}                   ; restore context
30      BX LR                          ; return
31
32  END

```

Figure 3: `asm_max.s`

Assembly code to find the maximum value in the array can be found in `asm_math.s`, while the C implementation as well as the CMSIS-DSP library implementation (using the `arm_max_f32()` function) is found in `main.c`. Take some time to study the code.

Figure 4 illustrates the expected output upon entering *Debug* mode and executing the code. We will now shift our attention to timing the code.

Figure 4: Expected output of the *max* example provided.

Timing

We can time the code by placing breakpoints around the sections of code that need to be timed and using the *t1* stopwatch provided by Keil. Figure 5 shows the breakpoints added around the three different implementations we wish to time, as well as the menu for the stopwatch *t1* and the option used to reset the time.

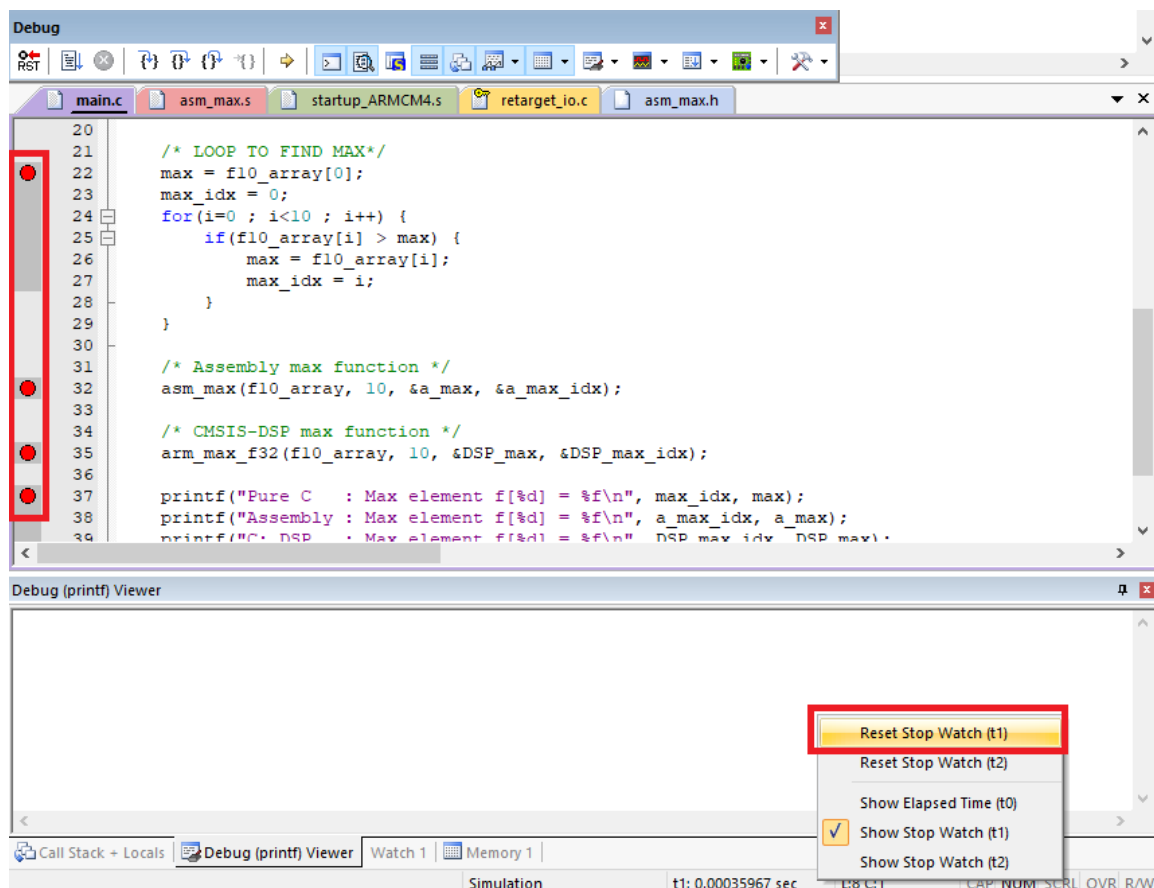


Figure 5: Adding breakpoints and the stopwatch menu.

1. We begin by executing the code until the first breakpoint and then resetting the stopwatch. We are now ready to time the C implementation.

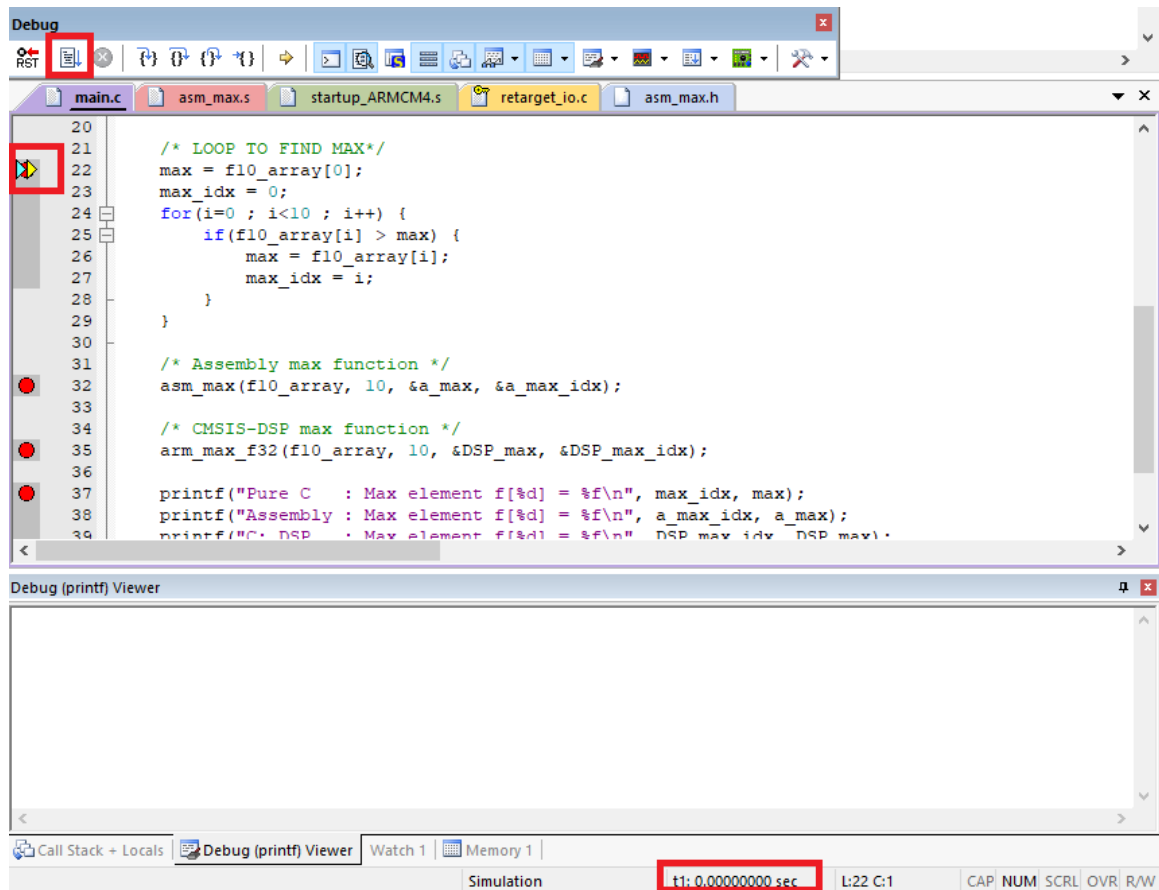


Figure 6: Run till first breakpoint and reset the stopwatch.

2. We run the code till the second breakpoint and note down the time taken, which is $14.33 \mu\text{s}$.

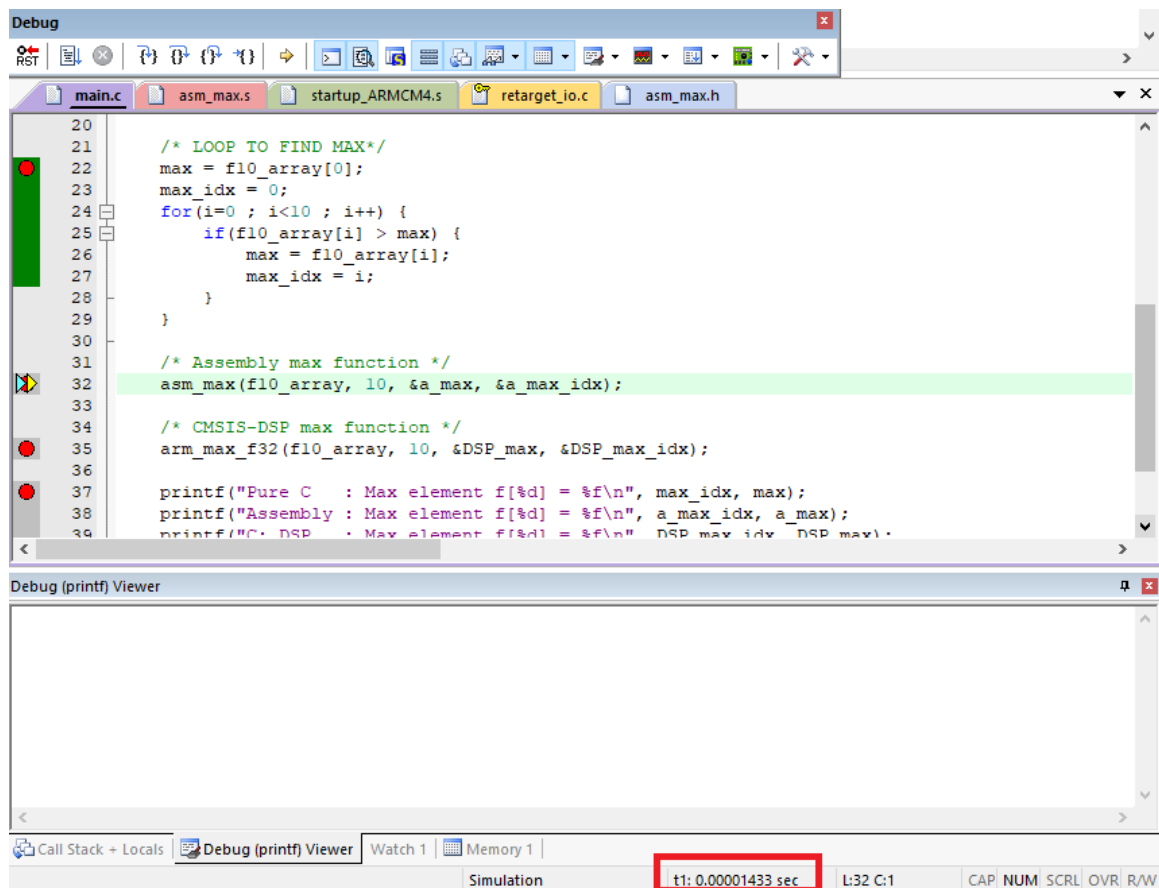


Figure 7: Run till second breakpoint and note down the time elapsed for the C implementation.

3. We now reset the stopwatch in order to time the assembly implementation.

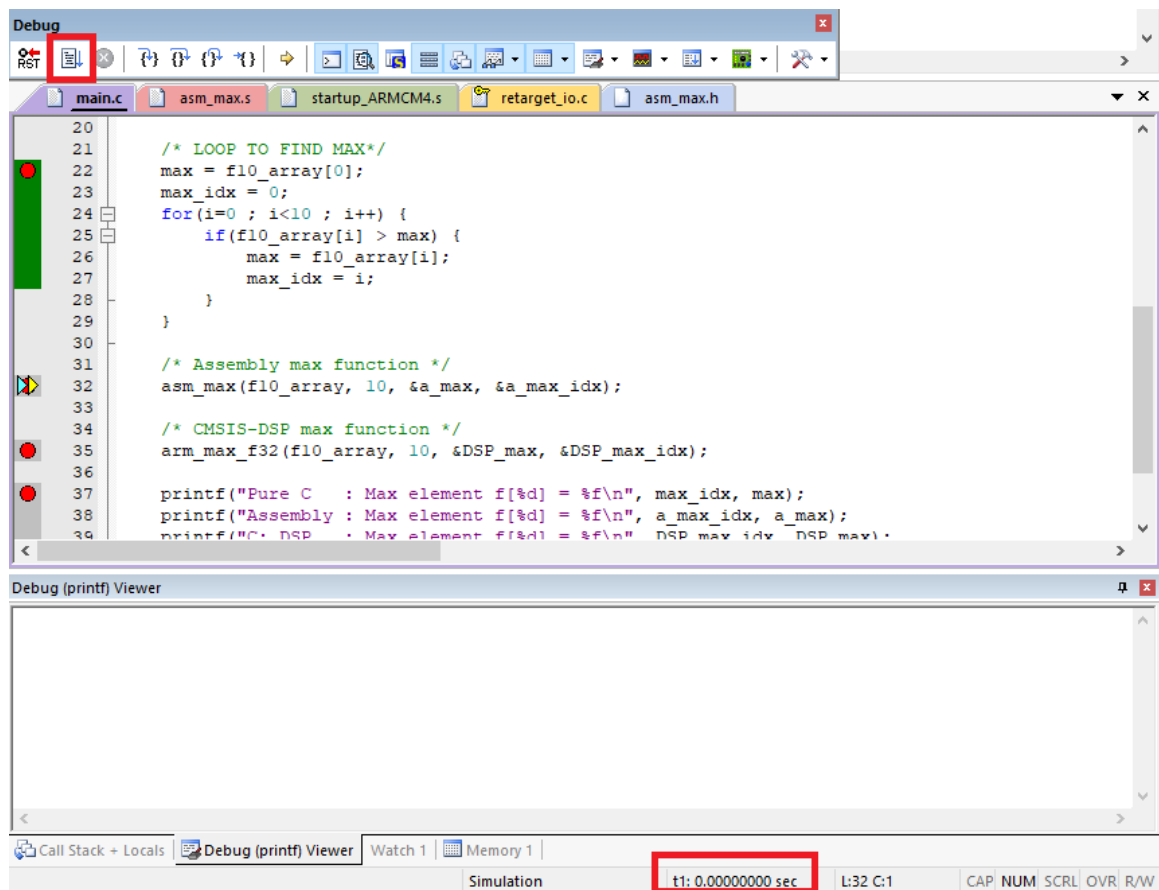


Figure 8: Reset the stopwatch.

4. We execute the code to the next breakpoint and note down the time taken, which is $13.17 \mu\text{s}$. Therefore the assembly implementation is roughly 8.81 % faster than the C implementation.

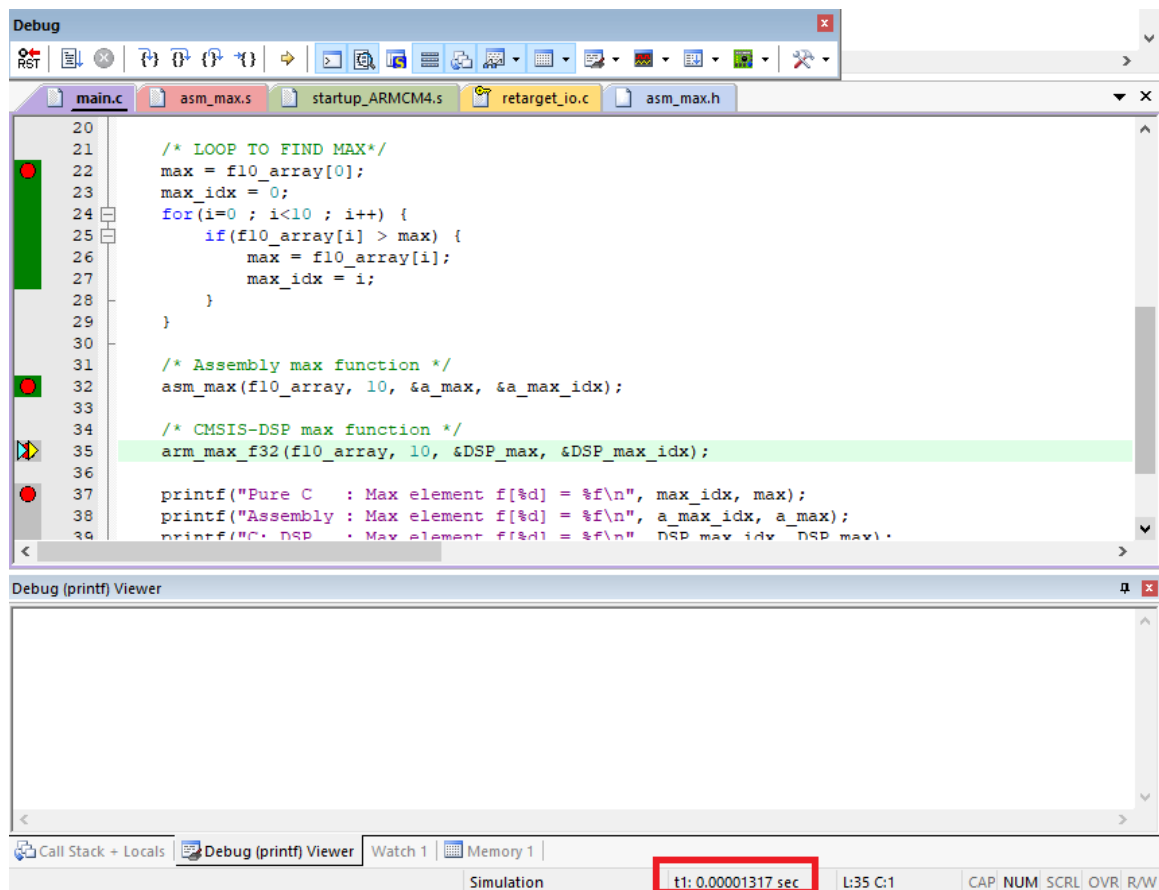


Figure 9: Run till third breakpoint and note down the time elapsed for the assembly implementation.

5. We now reset the stopwatch in order to time the CMSIS-DSP implementation.

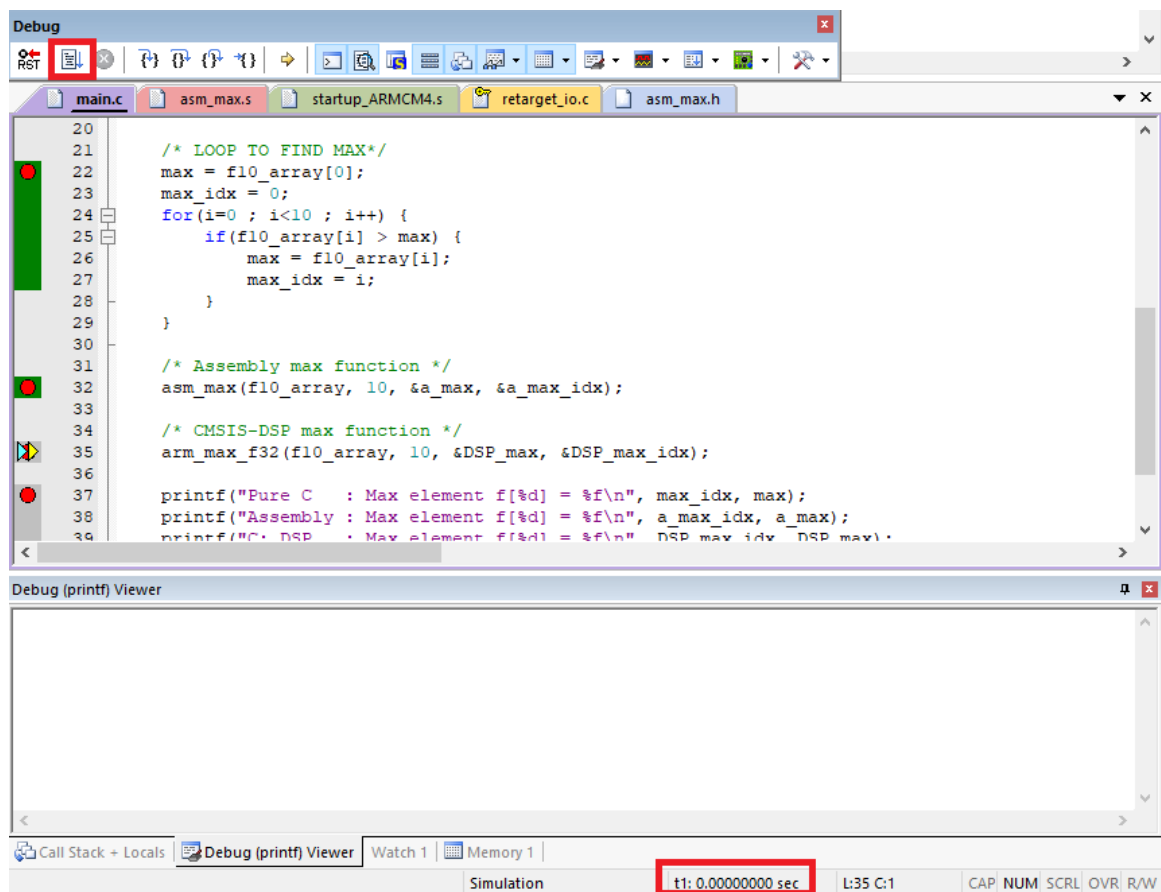


Figure 10: Reset the stopwatch.

6. We execute the code to the next breakpoint and note down the time taken, which is 9.92 μ s. The CMSIS-DSP implementation is 44.5% and 32.8% faster than the C and assembly implementation respectively. This simple example clearly illustrates the advantage of using the optimized CMSIS-DSP library. **Ensure that you can repeat this exercise using the base-project and source code provided.**

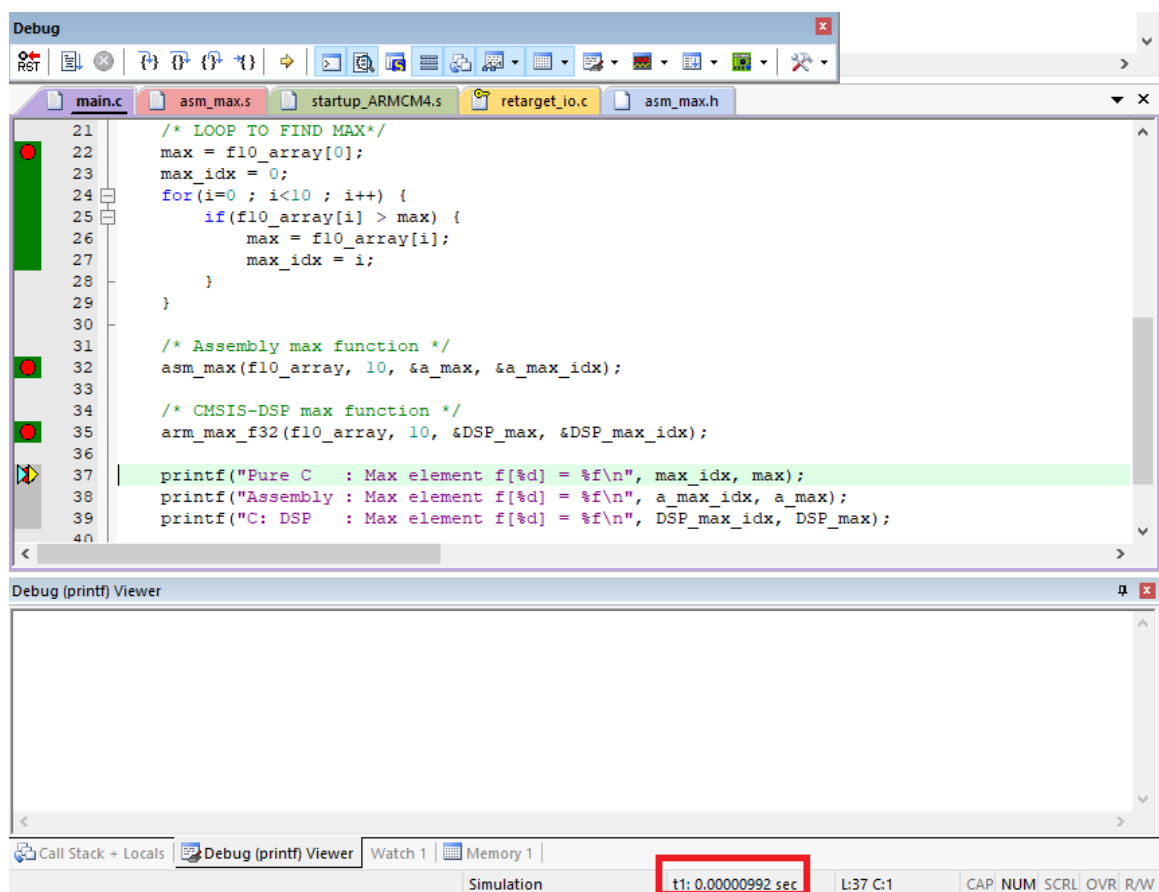


Figure 11: Run till final breakpoint and measure elapsed time for the CMSIS-DSP implementation.

Lab tasks

In this lab, you will complete two tasks based on the concepts introduced. You may use the same base-project provided to write code in, and refer to the [CMSIS-DSP documentation](#) to locate the library functions you may need.

Dot-product

Given two vectors (arrays) A and B each of length N , the dot-product DP can be computed by:

$$DP = \sum_{i=0}^{N-1} A(i) \times B(i)$$

- Write plain C code to calculate the dot-product. Then make use of a CMSIS-DSP library function to calculate the dot-product. Report the time taken by both implementations and the speed-up to your TA when presenting your code.
- Write an assembly implementation of the dot-product and measure the time taken.

Variance

Given a vector (array) A of length N , the variance σ^2 can be computed by:

$$\sigma^2 = \frac{\sum_{i=0}^{N-1} (A(i) - \mu)^2}{N}$$

where μ is the mean of vector A , computed by:

$$\mu = \frac{\sum_{i=0}^{N-1} A(i)}{N}$$

- Write plain C code to calculate the variance. Then make use of a CMSIS-DSP library function to calculate the variance. Report the time taken by both implementations and the speed-up to your TA when presenting your code.
- Write an assembly implementation of the variance and measure the time taken.

Assumptions

- All arrays and results are 32-bit floating point numbers.
- Use only the **f32** versions of CMSIS-DSP functions.
- Measure execution time for an array size of $N = 1000$ for both tasks.